

# Floating point numbers

Troels Henriksen

29-11-2021

Based on slides by Randal E. Bryant and David R. O'Hallaron.  
Some material by Michael Kirkedal Tomsen.

# Agenda

Preliminaries: biased numbers

Floating point arithmetic

- Background: Fractional binary numbers

- IEEE floating point standard

- Examples and properties

- Rounding, addition, and multiplication

- Floating point in C

Summary

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
00000000 <sub>2</sub>		

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
<hr/> 00000000 <sub>2</sub>	0 <sub>10</sub>	

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
00000000 <sub>2</sub>	0 <sub>10</sub>	-127 <sub>10</sub>
01111111 <sub>2</sub>		

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
00000000 <sub>2</sub>	0 <sub>10</sub>	-127 <sub>10</sub>
01111111 <sub>2</sub>	127 <sub>10</sub>	

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , **Bias** = 127

	B2U	B2I
00000000 <sub>2</sub>	0 <sub>10</sub>	-127 <sub>10</sub>
01111111 <sub>2</sub>	127 <sub>10</sub>	0 <sub>10</sub>
11111111 <sub>2</sub>		

# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
00000000 <sub>2</sub>	0 <sub>10</sub>	-127 <sub>10</sub>
01111111 <sub>2</sub>	127 <sub>10</sub>	0 <sub>10</sub>
11111111 <sub>2</sub>	255 <sub>10</sub>	



# Biased number representation

For *biased numbers*, the raw bits are interpreted as unsigned, and then a constant *bias* is subtracted.

**Unsigned**

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

**Two's complement**

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

**Biased**

$$B2I(X) = \left( \sum_{i=0}^{w-1} x_i \cdot 2^i \right) - \text{Bias}$$

- Typically

$$\text{Bias} = 2^{w-1} - 1$$

- Examples for  $w = 8$ , Bias = 127**

	B2U	B2I
00000000 <sub>2</sub>	0 <sub>10</sub>	-127 <sub>10</sub>
01111111 <sub>2</sub>	127 <sub>10</sub>	0 <sub>10</sub>
11111111 <sub>2</sub>	255 <sub>10</sub>	128 <sub>10</sub>

Preliminaries: biased numbers

Floating point arithmetic

- Background: Fractional binary numbers

- IEEE floating point standard

- Examples and properties

- Rounding, addition, and multiplication

- Floating point in C

Summary

# Integral binary numbers

We have seen that

$$10010101_2$$

is basically interpreted like

$$149_{10}$$

in particular “structure” is the same, just with 2 instead of 10.

# Integral binary numbers

We have seen that

$$10010101_2$$

is basically interpreted like

$$149_{10}$$

in particular “structure” is the same, just with 2 instead of 10.

**Can we do the same thing for fractional numbers?**

$$1011.101_2$$

## Fractional numbers

$$123.456 = 1 \cdot 10^2 + 2 \cdot 10^1 + 3 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2} + 6 \cdot 10^{-3}$$

# Fractional numbers

$$123.456 = 1 \cdot 10^2 + 2 \cdot 10^1 + 3 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2} + 6 \cdot 10^{-3}$$

Generally

$$a_{m-1} \cdots a_0 . a_{-1} \cdots a_{-n} = \sum_{i=-n}^{m-1} a_i \cdot 10^i$$

# Fractional numbers

$$123.456 = 1 \cdot 10^2 + 2 \cdot 10^1 + 3 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2} + 6 \cdot 10^{-3}$$

Generally

$$a_{m-1} \cdots a_0.a_{-1} \cdots a_{-n} = \sum_{i=-n}^{m-1} a_i \cdot 10^i$$

Even more generally, for radix  $r$

$$a_{m-1} \cdots a_0.a_{-1} \cdots a_{-n} = \sum_{i=-n}^{m-1} a_i \cdot r^i$$

# Fractional binary numbers

<b>Weight</b>	$2^{m-1}$	$2^{m-2}$	$\dots$	4	2	1		$1/2$	$1/4$	$1/8$	$\dots$	$2^{-n}$
<b>Digits</b>	$b_{m-1}$	$b_{m-2}$	$\dots$	$b_2$	$b_1$	$b_0$		$b_{-1}$	$b_{-2}$	$b_{-3}$	$\dots$	$b_{-n}$

## Representation

- Bits to the right of “binary point” represents fractional powers of 2.
- Represents rational number

$$b_{m-1} \dots b_0 . b_{-1} \dots b_{-n} = \sum_{i=-n}^{m-1} b_i \cdot 2^i$$



# Examples of fractional binary numbers

Value	Representation
$5\frac{3}{4}$	$101.11_2$
$2\frac{7}{8}$	$10.111_2$
$1\frac{7}{16}$	$1.0111_2$

## Observations

- Divide by 2 by logical shifting right.
- Multiply by 2 by shifting left.
- Numbers of form  $0.111\dots$  are just below 1.0.
  - ▶  $1/2 + 1/4 + 1/8 + \dots 1/2^n + \dots \sim 1.0$ .
  - ▶ Use notation  $1.0 - \epsilon$ .

# Representable numbers

## Limitation #1

- Can only represent fractional part of form  $x/2^k$
- Other rational numbers have repeating bit representation

Value	Representation
$\frac{1}{3}$	$0.0101010101[01] \cdots_2$
$\frac{1}{5}$	$0.001100110011[0011] \cdots_2$
$\frac{1}{10}$	$0.0001100110011[0011] \cdots_2$

## Limitation #2

- Just one setting of binary point within the  $w$  bits.
  - ▶ Limited range of numbers—very small values? Very large?

# The fixed-point dilemma

**Consider  $w = 8$**

## **1 bit for fraction**

- Largest number:  $1111111.1_2 = 127.5_{10}$
- Increment:  $0000000.1_2 = 0.5_{10}$

## **7 bits for fraction**

- Largest number:  $1.1111111_2 = 1.9921875_{10}$
- Increment:  $0.0000001_2 = 0.0078125_{10}$

## **4 bits for fraction**

- Largest number:  $1111.1111_2 = 15.9375_{10}$
- Increment:  $0000.0001_2 = 0.0625_{10}$

**Fixed-point has same absolute precision everywhere, but this means relative precision is worse for numbers close to 0!**

Preliminaries: biased numbers

Floating point arithmetic

Background: Fractional binary numbers

IEEE floating point standard

Examples and properties

Rounding, addition, and multiplication

Floating point in C

Summary

# IEEE Floating Point

## IEEE Standard 754

- Established in 1985 as uniform standard for floating point.
  - ▶ Many idiosyncratic formats before then.
- Supported by all major CPUs.

## Driven by numerical concerns

- Nice standards for rounding, overflow, underflow.
- Hard to make fast in hardware.
  - ▶ Numerical analysts predominated over hardware designers in defining standard.
  - ▶ ... but (later) Turing Award winner William Kahan secretly knew that Intel had figured out how.
  - ▶ **Beware the wrath of Kahan!**
  - ▶ <http://people.eecs.berkeley.edu/~wkahan/>

# Floating Point Representation


## Numerical form

$$(-1)^S \cdot M \cdot 2^E$$

- **Sign bit**  $S$  determines whether number negative or positive.
- **Significand**  $M$  normally a fractional value in range  $[1, 2)$ .
- **Exponent**  $E$  weights value by power of two.

## Encoding

- Most significant bit is sign bit.
- Exp field encodes  $E$  (but is not equal to  $E$ ).
- Frac field encodes  $M$  (but is not equal to  $M$ ).




# Why such a weird format?

## The point is floating

- No fixed number of bits allocated to “fraction”.
- More bits close to 0, fewer bits for numbers with large magnitude.
- Symmetric around 0.

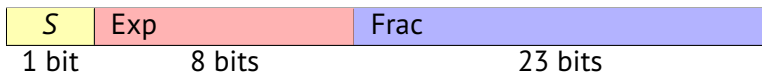
## Density of floats



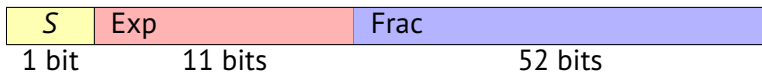
<https://stackoverflow.com/a/24179424/6131552>

# Precision options


## 32-bit single precision: float



## 64-bit double precision: double




## 80-bit Extended precision (Intel only, never use): long double





## Normalised values when $\text{Exp} \neq 0 \dots 0$ and $\text{Exp} \neq 1 \dots 1$

$$v = (-1)^S \cdot M \cdot 2^E$$



### ■ Exponent encoded as *biased* value

$$E = \text{Exp} - \text{Bias}$$

- ▶ Exp: unsigned value of Exp field.
- ▶ Bias =  $2^{k-1} - 1$ , where  $k$  is number of Exp bits.
  - ▶ Single precision: 127 ( $\text{Exp} \in [1, 254]$ ,  $E \in [-126, 127]$ ).
  - ▶ Double precision: 1023 ( $\text{Exp} \in [1, 2046]$ ,  $E \in [-1022, 1023]$ ).

### ■ Significand coded with implied leading 1:

$$M = 1.\text{xxx} \dots \text{x}_2$$

- ▶  $\text{xxx} \dots \text{x}$ : bits of Frac field.
- ▶ Minimum when Frac = 0000  $\dots$  0 ( $M = 1$ ).
- ▶ Maximum when Frac = 1111  $\dots$  1 ( $M = 2 - \epsilon$ ).
- ▶ Get extra leading bit for free.

# Normalised encoding example

$$v = (-1)^S \cdot M \cdot 2^E$$

$$E = \text{Exp} - \text{Bias}$$

Value: float F = 15213.0

$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \cdot 2^{13} \end{aligned}$$

Significand

$$\begin{aligned} M &= 1.1101101101101_2 \\ \text{Frac} &= 11011011011010000000000_2 \end{aligned}$$

Exponent

$$\begin{aligned} E &= 13_{10} \\ \text{Bias} &= 127_{10} \\ \text{Exp} &= 140_{10} = 10001100_2 \end{aligned}$$

Result	0	10001100	11011011011010000000000
--------	---	----------	-------------------------

# Denormal values

$$v = (-1)^S \cdot M \cdot 2^E$$

$$E = 1 - \text{Bias}$$

Occur when  $\text{Exp} = 000 \cdots 0_2$ .

- Exponent encoded as

$$E = 1 - \text{Bias}$$

- Significand coded with implied leading 0:

$$M = 0.xxx \cdots x_2$$

- Cases

- ▶  $\text{Exp} = 000 \cdots 0_2, \text{Frac} = 000 \cdots 0_2$ 
  - ▶ Represents zero value.
  - ▶ Note distinct values  $-0, +0$  – why do you think that is?
- ▶  $\text{Exp} = 000 \cdots 0_2, \text{Frac} \neq 000 \cdots 0_2$ 
  - ▶ Numbers closest to 0.0.
  - ▶ Called **Subnormal numbers**.
  - ▶ Ensure that  $x \neq y \Rightarrow x - y \neq 0$ , i.e. avoid overflow.

# Special values

**Occur when  $\text{Exp} = 111 \cdots 1_2$ .**

When  $\text{Exp} = 111 \cdots 1_2, \text{Frac} = 000 \cdots 0_2$

- Represents  $\pm\infty$ .
- Typically the result of *overflow*.
  - ▶ Overflow can be negative!
  - ▶ *Underflow* is when the result becomes zero due to rounding.
- Both positive and negative.
- Examples:

$$\frac{1}{0} = \frac{-1}{-0} = \infty \qquad \frac{1}{-0} = -\infty$$

When  $\text{Exp} = 111 \cdots 1_2, \text{Frac} \neq 000 \cdots 0_2$

- Not A Number (NaN).
- Represents case when no numeric value can be determined.
- Examples:

$$\text{sqrt}(-1) \qquad \infty - \infty \qquad \infty \cdot 0$$

# The floating point number line

← very positive E    very negative E →    ← very negative E    very positive E →

$-\infty$	-Normal	-Denorm	-0	+0	+Denorm	+Normal	$+\infty$
-----------	---------	---------	----	----	---------	---------	-----------

NaN
-----

NaN
-----

Note that NaNs are unordered:

- NaN is different from everything *even other NaNs!*
  - ▶ `NaN == NaN` is false.
  - ▶ Floating-point equality is not reflexive!
- `NaN > x` and `NaN < x` is false for all `x`.

Preliminaries: biased numbers

Floating point arithmetic

Background: Fractional binary numbers

IEEE floating point standard

Examples and properties

Rounding, addition, and multiplication

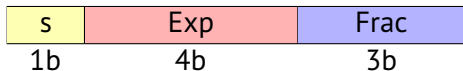
Floating point in C

Summary

Play the game

<https://topps.diku.dk/compsys/floating-point.html>

# Tiny 8-bit floating point example



## 8-bit floating point representation

- Sign bit is the most significant bit (leftmost).
- The next four bits are Exp with a bias of 7.
- The last three bits are Frac.

## Same general form as IEEE Format


- Normalised, denormalised.
- Representation of 0, NaN, both infinities.



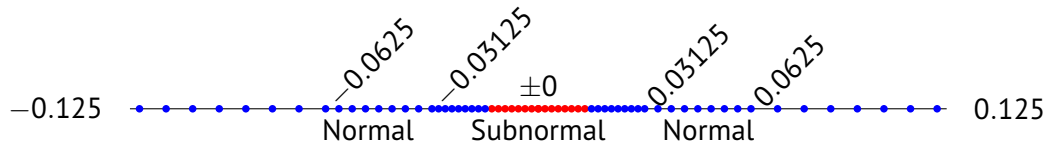
# Dynamic range of positive numbers

	Sign	Exp	Frac	E	Value	
Denormalised	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 \cdot 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 \cdot 1/64 = 2/512$	
	...					
	0	0000	111	-6	$7/8 \cdot 1/64 = 7/512$	largest denorm
Normalised	0	0001	000	-6	$8/8 \cdot 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 \cdot 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 \cdot 1/2 = 14/16$	
	0	0110	111	-1	$15/8 \cdot 1/2 = 15/16$	
	0	0111	000	0	$8/8 \cdot 1 = 1$	Closest to 1
	0	0111	001	0	$9/8 \cdot 9/8 = 1$	
	0	0111	010	0	$10/8 \cdot 10/8 = 1$	
	...					
	0	1110	110	7	$14/8 \cdot 128 = 224$	
	0	1110	111	7	$15/8 \cdot 128 = 240$	
	0	1111	000	N/A	$\infty$	

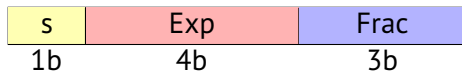
# Distribution of values



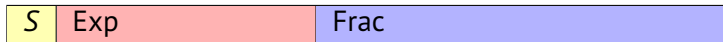
## Distribution of values (zooming in)



- Note how the distribution gets denser towards zero.
- Note the big gap there would be around 0 if we did not have subnormals.
- Each of the spans with same distance between neighbors corresponds to numbers with same Exp.



# Useful properties of the IEEE encoding



- **Floating-point zero same as integer zero**
  - ▶ All bits 0.
  - ▶ ...but negative zero is different.
- **Can almost compare floats with unsigned integer comparisons**
  - ▶ Must first compare sign bit.
  - ▶ Must consider  $-0 = 0$ .
  - ▶ NaNs problematic:
    - ▶ Greater than any other value (because  $\text{Exp} = 111 \cdots 1_2$ ).
    - ▶ What should comparison yield?
  - ▶ Otherwise OK:
    - ▶ Normalised and denormalised compare as expected.
    - ▶ Infinities ordered properly relative to finities.

Preliminaries: biased numbers

Floating point arithmetic

Background: Fractional binary numbers

IEEE floating point standard

Examples and properties

Rounding, addition, and multiplication

Floating point in C

Summary

# Basic idea behind floating point operations

$$x +_f y = \text{Round}(x + y)$$

$$x \times_f y = \text{Round}(x \times y)$$

- **Basic idea**

- ▶ First *compute exact result!*
- ▶ Then round it to fit into desired precision.
  - ▶ Overflow if exponent too large.
  - ▶ *Round to fit* into Frac.

# Rounding and rounding modes

- There's more than one way to round a number, here to an integer.

	1.40	1.60	1.50	2.50	-1.50
Towards zero					

# Rounding and rounding modes

- There's more than one way to round a number, here to an integer.

	1.40	1.60	1.50	2.50	-1.50
Towards zero	1	1	1	2	-1
Towards $-\infty$					



# Rounding and rounding modes

- There's more than one way to round a number, here to an integer.

	1.40	1.60	1.50	2.50	-1.50
Towards zero	1	1	1	2	-1
Towards $-\infty$	1	1	1	2	-2
Towards $\infty$					

# Rounding and rounding modes

- There's more than one way to round a number, here to an integer.

	1.40	1.60	1.50	2.50	-1.50
Towards zero	1	1	1	2	-1
Towards $-\infty$	1	1	1	2	-2
Towards $\infty$	2	2	2	3	-1
Nearest even					

# Rounding and rounding modes

- There's more than one way to round a number, here to an integer.

	1.40	1.60	1.50	2.50	-1.50
Towards zero	1	1	1	2	-1
Towards $-\infty$	1	1	1	2	-2
Towards $\infty$	2	2	2	3	-1
Nearest even $\infty$	1	2	2	2	-2

- "Round to nearest, ties to even" is the default rounding mode.

## Closer look at *nearest even*

- **Default rounding mode**

- ▶ But can be changed dynamically.
  - ▶ `https://www.gnu.org/software/libc/manual/html_node/Rounding.html`
  - ▶ Never do this.
- ▶ All others are statistically biased.
  - ▶ Biased: Sum of set of positive numbers will consistently be over- or under-estimated.

- **Applying to other decimal places / bit positions**

- ▶ When exactly halfway between two possible values:
  - ▶ Round so that least significant digit is even.
- ▶ E.g. rounding to nearest hundredth:
  - ▶ 7.8949999:

## Closer look at *nearest even*

- **Default rounding mode**

- ▶ But can be changed dynamically.
  - ▶ `https://www.gnu.org/software/libc/manual/html_node/Rounding.html`
  - ▶ Never do this.
- ▶ All others are statistically biased.
  - ▶ Biased: Sum of set of positive numbers will consistently be over- or under-estimated.

- **Applying to other decimal places / bit positions**

- ▶ When exactly halfway between two possible values:
  - ▶ Round so that least significant digit is even.
- ▶ E.g. rounding to nearest hundredth:
  - ▶ 7.8949999: 7.89
  - ▶ 7.8990001:

## Closer look at *nearest even*

- **Default rounding mode**

- ▶ But can be changed dynamically.
  - ▶ `https://www.gnu.org/software/libc/manual/html_node/Rounding.html`
  - ▶ Never do this.
- ▶ All others are statistically biased.
  - ▶ Biased: Sum of set of positive numbers will consistently be over- or under-estimated.

- **Applying to other decimal places / bit positions**

- ▶ When exactly halfway between two possible values:
  - ▶ Round so that least significant digit is even.
- ▶ E.g. rounding to nearest hundredth:
  - ▶ 7.8949999: 7.89
  - ▶ 7.8990001: 7.90
  - ▶ 7.8950000:

## Closer look at *nearest even*

- **Default rounding mode**

- ▶ But can be changed dynamically.
  - ▶ [https://www.gnu.org/software/libc/manual/html\\_node/Rounding.html](https://www.gnu.org/software/libc/manual/html_node/Rounding.html)
  - ▶ Never do this.
- ▶ All others are statistically biased.
  - ▶ Biased: Sum of set of positive numbers will consistently be over- or under-estimated.

- **Applying to other decimal places / bit positions**

- ▶ When exactly halfway between two possible values:
  - ▶ Round so that least significant digit is even.
- ▶ E.g. rounding to nearest hundredth:
  - ▶ 7.8949999: 7.89
  - ▶ 7.8990001: 7.90
  - ▶ 7.8950000: 7.90
  - ▶ 7.8850000:

## Closer look at *nearest even*

- **Default rounding mode**

- ▶ But can be changed dynamically.
  - ▶ [https://www.gnu.org/software/libc/manual/html\\_node/Rounding.html](https://www.gnu.org/software/libc/manual/html_node/Rounding.html)
  - ▶ Never do this.
- ▶ All others are statistically biased.
  - ▶ Biased: Sum of set of positive numbers will consistently be over- or under-estimated.

- **Applying to other decimal places / bit positions**

- ▶ When exactly halfway between two possible values:
  - ▶ Round so that least significant digit is even.
- ▶ E.g. rounding to nearest hundredth:
  - ▶ 7.8949999: 7.89
  - ▶ 7.8990001: 7.90
  - ▶ 7.8950000: 7.90
  - ▶ 7.8850000: 7.88



# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
-------	--------	---------	--------	---------------

---

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$				

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00011_2$			

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{011}_2$	$10.00_2$	( $< 1/2$ – down)	2

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{011}_2$	$10.00_2$	( $< 1/2$ – down)	2
$2 \frac{3}{16}$				

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00011_2$	$10.00_2$	( $< 1/2$ – down)	2
$2 \frac{3}{16}$	$10.00110_2$			

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ –down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{11}0_2$	$10.01_2$	( $> 1/2$ –up)	$2 \frac{1}{4}$

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ –down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{11}0_2$	$10.01_2$	( $> 1/2$ –up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$				



# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ –down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	$10.01_2$	( $> 1/2$ –up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$			

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ -down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	$10.01_2$	( $> 1/2$ -up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	$11.00_2$	( $1/2$ -up)	3

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \cdots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ –down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	$10.01_2$	( $> 1/2$ –up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	$11.00_2$	( $1/2$ –up)	3
$2 \frac{5}{8}$				

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ -down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	$10.01_2$	( $> 1/2$ -up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	$11.00_2$	( $1/2$ -up)	3
$2 \frac{5}{8}$	$10.10\textcolor{red}{100}_2$			

# Rounding binary numbers

- **Binary fractional numbers**

- ▶ “Even” when least significant bit is 0.
- ▶ “Half way” when bits to right of rounding position are  $100 \dots_2$ .

- **Examples**

- ▶ Round to nearest  $1/4$  (2 bits right of binary point).

Value	Binary	Rounded	Action	Rounded value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	$10.00_2$	( $< 1/2$ –down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	$10.01_2$	( $> 1/2$ –up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	$11.00_2$	( $1/2$ –up)	3
$2 \frac{5}{8}$	$10.10\textcolor{red}{100}_2$	$10.10_2$	( $1/2$ –down)	$2 \frac{1}{2}$

# Floating point multiplication (assuming operands are numbers)

$$((-1)^{S_3} \cdot M_3 \cdot 2^{E_3}) = ((-1)^{S_1} \cdot M_1 \cdot 2^{E_1}) \cdot ((-1)^{S_2} \cdot M_2 \cdot 2^{E_2})$$

- **Exact result**

$$S_3 = S_1 \oplus S_2$$

$$M_3 = M_1 \cdot M_2$$

$$E_3 = E_1 + E_2$$

where  $\oplus$  is exclusive-or.

- **Fixing**

- ▶ If  $M_3 \geq 2$ , shift  $M_3$  right and increment  $E_e$ .
- ▶ If  $E_3$  out of range, overflow to  $\infty$ .
- ▶ Round  $M_3$  to fit Frac precision.

- **Implementation**

- ▶ Biggest chore is multiplying significands.

# Floating point addition (assuming operands are numbers)

$$((-1)^{S_3} \cdot M_3 \cdot 2^{E_3}) = ((-1)^{S_1} \cdot M_1 \cdot 2^{E_1}) + ((-1)^{S_2} \cdot M_2 \cdot 2^{E_2})$$

## ■ Approach

- ▶ Assume without loss of generality that  $E_1 \geq E_2$ .
- ▶ Rewrite smaller number such that its exponent matches  $E_1$ :

$$((-1)^{S_3} \cdot M_3 \cdot 2^{E_3}) = ((-1)^{S_1} \cdot M_1 \cdot 2^{E_1}) + ((-1)^{S_2} \cdot M'_2 \cdot 2^{E_1})$$

## ■ Exact result

- ▶ Sign  $S_3$ , significant  $M_3$ :
  - ▶ Result of signed addition.

## ■ Fixing

- ▶ If  $M_3 \geq 2$ , shift  $M_3$  right and increment  $E_3$ .
- ▶ If  $M_3 < 1$ , shift  $M$  left  $k$  positions and decrement  $E_3$  by  $k$ .
- ▶ If  $E_3$  out of range, overflow to  $\infty$ .
- ▶ Round  $M$  to fit Frac precision.

$$\begin{array}{r} \leftarrow E_1 - E_2 \rightarrow \\ \boxed{-1^{S_1} \cdot M_1} \\ + \quad \boxed{-1^{S_2} \cdot M_2} \\ \hline \boxed{-1^{S_3} \cdot M_3} \end{array}$$

## Example of floating-point addition with a 2-bit significand

$$\begin{aligned} & (-1.01 \cdot 2^2) + (1.1 \cdot 2^4) \\ = & (-1.01 \cdot 2^2) + (110.0 \cdot 2^2) && \text{Align exponents} \\ = & (-1.01 + 110.0) \cdot 2^2 && \text{Distributivity} \\ = & 100.11 \cdot 2^2 && \text{Add significands} \\ = & 1.0011 \cdot 2^4 && \text{Normalise} \\ = & 1.01 \cdot 2^4 && \text{Perform rounding} \end{aligned}$$



# Algebraic properties of floating-point addition

- **Compared to those of Abelian Group**
  - ▶ Closed under addition?

# Algebraic properties of floating-point addition

- **Compared to those of Abelian Group**

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative?

# Algebraic properties of floating-point addition

- **Compared to those of Abelian Group**

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative?

# Algebraic properties of floating-point addition

- **Compared to those of Abelian Group**

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(3.14 + 1e10) - 1e10 = 0$
  - ▶  $3.14 + (1e10 - 1e10) = 3.14$
- ▶ 0 is additive identity?

# Algebraic properties of floating-point addition

- **Compared to those of Abelian Group**

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(3.14 + 1e10) - 1e10 = 0$
  - ▶  $3.14 + (1e10 - 1e10) = 3.14$
- ▶ 0 is additive identity? **Yes**
- ▶ Does every element have an additive inverse?

# Algebraic properties of floating-point addition

## ■ Compared to those of Abelian Group

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(3.14 + 1e10) - 1e10 = 0$
  - ▶  $3.14 + (1e10 - 1e10) = 3.14$
- ▶ 0 is additive identity? **Yes**
- ▶ Does every element have an additive inverse? **Almost**
  - ▶ Infinities and NaN do not have inverses.

## ■ Monotonicity

- ▶  $a \geq b \Rightarrow a + c \geq b + c?$

# Algebraic properties of floating-point addition

## ■ Compared to those of Abelian Group

- ▶ Closed under addition? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(3.14 + 1e10) - 1e10 = 0$
  - ▶  $3.14 + (1e10 - 1e10) = 3.14$
- ▶ 0 is additive identity? **Yes**
- ▶ Does every element have an additive inverse? **Almost**
  - ▶ Infinities and NaN do not have inverses.

## ■ Monotonicity

- ▶  $a \geq b \Rightarrow a + c \geq b + c$ ? **Almost**
  - ▶ Infinities and NaNs are the exception.

# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**
  - ▶ Closed under multiplication?



# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**
  - ▶ Closed under multiplication? **Yes**
    - ▶ But may generate infinity or NaN.
  - ▶ Commutative?

# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**

- ▶ Closed under multiplication? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative?

# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**

- ▶ Closed under multiplication? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(1e20 * 1e20) * 1e-20 = \infty$
  - ▶  $1e20 * (1e20 * 1e-20) = 1e20$
- ▶ 1 is multiplicative identity?

# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**

- ▶ Closed under multiplication? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(1e20 * 1e20) * 1e-20 = \infty$
  - ▶  $1e20 * (1e20 * 1e-20) = 1e20$
- ▶ 1 is multiplicative identity? **Yes**
- ▶ Multiplication distributes over addition?

# Algebraic properties of floating-point multiplication

- **Compared to those of a commutative ring**

- ▶ Closed under multiplication? **Yes**
  - ▶ But may generate infinity or NaN.
- ▶ Commutative? **Yes**
- ▶ Associative? **No**
  - ▶ Due to overflow and inexactness of rounding.
  - ▶  $(1e20 * 1e20) * 1e-20 = \infty$
  - ▶  $1e20 * (1e20 * 1e-20) = 1e20$
- ▶ 1 is multiplicative identity? **Yes**
- ▶ Multiplication distributes over addition? **No**
  - ▶ Overflow and rounding again.
  - ▶  $1e20 * (1e20 - 1e20) = 0.0$
  - ▶  $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$

Preliminaries: biased numbers

Floating point arithmetic

Background: Fractional binary numbers

IEEE floating point standard

Examples and properties

Rounding, addition, and multiplication

Floating point in C

Summary

# Floating point in C

- **C guarantees two types**

- ▶ `float`: 32-bit single precision.
- ▶ `double`: 64-bit single precision.

- **Conversions/casting**

- ▶ Casting between `int`, `float`, and `double` changes bit representation.
- ▶ `double/float to int`
  - ▶ Truncates fractional part.
  - ▶ Like rounding toward zero.
  - ▶ Not defined when out of range or NaN: generally sets to TMin.
- ▶ `int to double`
  - ▶ Exact conversion as long as `int` fits in 53 bits.
- ▶ `int to float`
  - ▶ Will round according to rounding mode.

Floating point is exciting!



**First “flight” of the Ariane 5 in 1996.**



# Floating point is exciting!



**First “flight” of the Ariane 5 in 1996.**

- A double storing horizontal velocity of the rocket was converted to a 16-bit signed integer.
- The number was larger than 32767 so the conversion failed, causing an exception, crashing the guidance module.

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

▪ `x == (int) (float) x`

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume `int` is 32 bits.

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

- `x == (int) (float) x`

- `x == (int) (double) x`

```
int    x = ...;
```

```
float  f = ...;
```

```
double d = ...;
```

Assume neither `d` nor `t`  
is NaN.

Assume `int` is 32 bits.

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`

# Floating point puzzles

**For each of the following C expressions, either**

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`

# Floating point puzzles

For each of the following C expressions, either

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`
- `2/3 == 2/3.0`



# Floating point puzzles

For each of the following C expressions, either

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor f  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`
- `2/3 == 2/3.0`
- `d < 0.0  $\Rightarrow$  (d*2) < 0.0`

# Floating point puzzles

For each of the following C expressions, either

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor f  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`
- `2/3 == 2/3.0`
- `d < 0.0  $\Rightarrow$  (d*2) < 0.0`
- `d > f  $\Rightarrow$  -f > -d`

# Floating point puzzles

For each of the following C expressions, either

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;
float  f = ...;
double d = ...;
```

Assume neither d nor t  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`
- `2/3 == 2/3.0`
- `d < 0.0  $\Rightarrow$  (d*2) < 0.0`
- `d > f  $\Rightarrow$  -f > -d`
- `d * d >= 0.0`

# Floating point puzzles

For each of the following C expressions, either

- Argue that it is true for all argument values.
- Explain why it's not.

```
int    x = ...;  
float  f = ...;  
double d = ...;
```

Assume neither d nor f  
is NaN.

Assume int is 32 bits.

- `x == (int) (float) x`
- `x == (int) (double) x`
- `f == (float) (double) f`
- `d == (double) (float) d`
- `f == -(-f)`
- `2/3 == 2/3.0`
- `d < 0.0  $\Rightarrow$  (d*2) < 0.0`
- `d > f  $\Rightarrow$  -f > -d`
- `d * d >= 0.0`
- `(d+f)-d == f`

Preliminaries: biased numbers

Floating point arithmetic

- Background: Fractional binary numbers

- IEEE floating point standard

- Examples and properties

- Rounding, addition, and multiplication

- Floating point in C

Summary

# Summary

- **IEEE floating point has clear properties.**
  - ▶ But they may not match your intuition.
- **Represents numbers of the form  $M \cdot 2^E$ .**
- One can reason about operations independent of implementation.
  - ▶ Computed with perfect precision and then rounded.
  - ▶ But rounded after *every* “primitive” operation (e.g. addition, multiplication).
- **Not the same as  $\mathbb{Q}/\mathbb{R}$  arithmetic.**
  - ▶ Violates associativity and distributivity, mostly due to rounding.
  - ▶ Sometimes makes life difficult for heavy-duty numerical programming.
  - ▶ But carefully designed such that “naive” use mostly does what one expects.