

# Byte-oriented memory, pointers, and IO

Troels Henriksen

02-12-2021


Based on slides by Randal E. Bryant and David R. O'Hallaron

A machine view of memory

A machine view of text

Binary IO

# Byte-oriented memory organisation



- **Programs refer to data by address**
  - ▶ Conceptually, envision as large array of bytes.
    - ▶ It's not really, but it works as a model.
  - ▶ An address is like an index into that array.
    - ▶ A *pointer* stores an address.
    - ▶ **Addresses are ultimately just unsigned integers.**
- **System provides private address space to each “process”**

## Any given computer has a “word size”

- “Native” size of integer-valued data.
  - ▶ But especially of addresses.
- 32-bit machines used to be the norm and are still found.
  - ▶  $2^{32}$  different addresses, meaning 4GiB can be addressed.
- 64-bit machines are most common.
  - ▶  $2^{64}$  different addresses, meaning 18EiB can be addressed.
  - ▶  $18.4 \cdot 10^{18}$  bytes.
  - ▶ Current machines only use lower 48 bits of address.
- Machines also support other data formats.
  - ▶ Fractions or multiples of word size.
  - ▶ Always integral number of types.
  - ▶ Smaller types (e.g. 16-bit integers) take less space in memory, but are (usually) not faster than the “native” words.
  - ▶ But bigger types (e.g. 128-bit integers) are slower.

# Word-oriented memory organisation

- **Addresses specify byte locations**

- ▶ Address of first byte in word.
- ▶ Addresses of successive words differ by 4 (32 bit) or 8 (64 bit).
- ▶ *Addresses always refer to a byte* even when addressing larger types.

- **We can take the address of any variable in a C program**

- ▶ `&x` gives us the address of `x`.
- ▶ If `x` has type `T`, then `&x` has type `T*`.

## Example data representations

<b>C type</b>	<b>Size in bytes on x86-64</b>
char	1
short	2
int	4
long	8
pointer	8

# Byte ordering

- **So, how are the bytes within a multi-byte word ordered in memory?**
  - ▶ Most significant byte at lowest address, or least significant byte at lowest address?
- **Conventions**
  - ▶ Big endian: SPARC, POWER, Internet protocols.
    - ▶ Least significant byte has highest address (“comes last”).
  - ▶ Little endian: x86, ARM (mostly).
  - ▶ Least significant byte has highest address (“comes first”).

# Byte ordering example

## ■ Example

- ▶ Variable has 4-byte value of 0x01234567.
- ▶ Address &x is 0x100.
  - ▶ No matter what, the address of an object is always the address of the *first* byte in the object (counting from lowest addresses).

Big endian

		0x100	0x101	0x102	0x103		
		01	23	45	67		

Little endian

		0x100	0x101	0x102	0x103		
		67	45	23	01		



# Byte ordering example

## ■ Example

- ▶ Variable has 4-byte value of 0x01234567.
- ▶ Address &x is 0x100.
  - ▶ No matter what, the address of an object is always the address of the *first* byte in the object (counting from lowest addresses).

### Big endian

		0x100	0x101	0x102	0x103		
		01	23	45	67		

### Little endian

		0x100	0x101	0x102	0x103		
		67	45	23	01		

## Important note

This difference is *not visible* unless you start decomposing integers as bytes with low-level operations. Bit-shifting etc. always acts as expected.

# Examining data representations

- **Code to print byte representation of data**

- ▶ Casting pointer to unsigned char\* allows treatment as byte array.

```
void show_bytes(unsigned char* start, size_t len) {  
    size_t i;  
    for (i = 0; i < len; i++)  
        printf("%p\t0x%.2x\n", start+i, start[i]);  
    printf("\n");  
}
```

## **printf directives:**

- %p: Print pointer.
- %x: Print hexadecimal.

## show\_bytes execution example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((unsigned char*) &a, sizeof(int));
```

### Result (Linux x86-64):

```
0x7fffb7f71dbc 6d  
0x7fffb7f71dbd 3b  
0x7fffb7f71dbe 00  
0x7fffb7f71dbf 00
```

A machine view of memory

A machine view of text

Binary IO

## Text IO

```
printf("Hello , world!\n");
```

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123



## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer: 123

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer:     123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;  
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer:     123

```
double y = 1.23;  
printf("a float: %f\n", y);
```

a float: 1.230000

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;  
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer:     123

```
double y = 1.23;  
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;  
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer:     123

```
double y = 1.23;  
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

a mess: 4202562

## Text IO

```
printf("Hello , world!\n");
```

Hello , world!

```
int x = 123;
```

```
printf("an integer: %d\n", x);
```

an integer: 123

```
printf("an integer: %5d\n", x);
```

an integer:     123

```
double y = 1.23;
```

```
printf("a float: %f\n", y);
```

a float: 1.230000

```
printf("a mess: %d\n", y);
```

a mess: 4202562

**Make sure format specifiers and argument types match!**

# Text representation

- **Machines only understand numbers, and text is an abstraction!**
- E.g. when the terminal receives a byte with the value 65, it draws an A.
- `printf()` determines which *bytes* must be written to the terminal to produce the text corresponding to e.g. the number 123: [49, 50, 51].

## Character sets

A character set maps a *number* to a *character*.

- ASCII defines characters in the range 0–127 ([asciitable.com](http://asciitable.com)).
- Some are invisible/unprintable *control characters*
- *Unicode* is a superset of ASCII that defines tens of thousands of characters for all the world's scripts.

**We'll assume ASCII, which has the simple property that 1 byte = 1 character.**



# The ASCII table

Control characters				Normal characters											
000	nul	016	dle	032	␣	048	0	064	@	080	P	096	'	112	p
001	soh	017	dc1	033	!	049	1	065	A	081	Q	097	a	113	q
002	stx	018	dc2	034	"	050	2	066	B	082	R	098	b	114	r
003	etx	019	dc3	035	#	051	3	067	C	083	S	099	c	115	s
004	eot	020	dc4	036	\$	052	4	068	D	084	T	100	d	116	t
005	enq	021	nak	037	%	053	5	069	E	085	U	101	e	117	u
006	ack	022	syn	038	&	054	6	070	F	086	V	102	f	118	v
007	bel	023	etb	039	'	055	7	071	G	087	W	103	g	119	w
008	bs	024	can	040	(	056	8	072	H	088	X	104	h	120	x
009	tab	025	em	041	)	057	9	073	I	089	Y	105	i	121	y
010	lf	026	eof	042	*	058	:	074	J	090	Z	106	j	122	z
011	vt	027	esc	043	+	059	;	075	K	091	[	107	k	123	{
012	np	028	fs	044	,	060	<	076	L	092	␣	108	l	124	
013	cr	029	gs	045	-	061	=	077	M	093	]	109	m	125	}
014	so	030	rs	046	.	062	>	078	N	094	^	110	n	126	~
015	si	031	us	047	/	063	?	079	O	095	_	111	o	127	del

## Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

## Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

The text *string* that is passed to `printf()` looks like this in memory:

Characters	x	:		%	d	\n	\0
Bytes	120	58	32	37	100	10	0

## Turning numbers into text

```
int x = 1234;  
printf("x: %d\n", x);
```

The text *string* that is passed to `printf()` looks like this in memory:

Characters	x	:		%	d	\n	\0
Bytes	120	58	32	37	100	10	0

`printf()` rewrites format specifiers (%d) to the textual representation of their corresponding value argument:

Characters	x	:		1	2	3	4	\n	\0
Bytes	120	58	32	49	50	51	52	10	0

These bytes (except the 0) are then written to *standard output* (typically the terminal) which interprets them as characters and eventually draws pixels on the screen.

# Machine representation versus text representation

```
int x = 305419896;
```

- Written as hexadecimal (base-16), this number is 0x12345678.
- One hexadecimal digit is 4 bit, so each group of two digits is one byte, and the number takes four bytes (32 bits).
- The *machine representation* in memory on an x86 CPU is  
0x78 0x56 0x34 0x12
- A *decimal text representation* in memory on *any* CPU is  
0x33 0x30 0x35 0x34 0x35 0x36 0x37 0x38
- Endianness has *no effect on text* (at least not with single-byte characters).
- In C, we have the additional convention that any string must be NUL-terminated.
- We identify a string with the address of its first character.

A machine view of memory

A machine view of text

Binary IO

## Writing bytes

The `fwrite` function writes raw data to an open file:

```
size_t fwrite(const void *ptr ,  
              size_t size ,  
              size_t nmemb,  
              FILE *stream );
```

`ptr`: the address in memory of the data.

`size`: the size of each data element in bytes.

`nmemb`: the number of data elements.

`stream`: the target file (opened with `fopen()`).

- Returns the number of data elements written (equal to `nmemb` unless an error occurs).
- Usually no difference between writing one `size * y` element or `x * size-y` elements—do whatever is convenient.

## Example of fwrite()

```
#include <stdio.h>

int main() {
    // Open for writing ("w")
    FILE *f = fopen("output", "w");

    char c = 42;

    fwrite(&c, sizeof(char), 1, f);

    fclose(f);
}
```

- Produces a file output.
- File contains the byte 42, corresponding to the ASCII character \*.
- **char is just an 8-bit integer type!**
  - ▶ No special “character” meaning.
  - ▶ Most Unicode characters will not fit in a single char (e.g. 'æ' needs 16 bits in UTF-8).
  - ▶ Name is unfortunate/historical.
  - ▶ Signedness is *implementation-defined* for historical reasons.



## Another example

```
#include <stdio.h>

int main() {
    FILE *f = fopen("output", "w");

    int x = 0x53505048;
    // Stored as 0x48 0x50 0x50 0x53

    fwrite(&x, sizeof(int), 1, f);

    fclose(f);
}
```

- Writes bytes 0x48 0x50 0x50 0x53.
- Corresponds to ASCII characters HPPS.
- A big-endian machine would produce SPPH.
- **Don't write code that depends on this!**

## Converting a non-negative integer to its ASCII representation

```
FILE *f = fopen("output", "w");
int x = 1337;           // Number to write;
char s[10];             // Output buffer.
int i = 10;             // Index of last character written.
while (1) {
    int d = x % 10;      // Pick out last decimal digit.
    x = x / 10;          // Remove last digit.
    i = i - 1;           // Index of next character.
    s[i] = '0' + d;      // Save ASCII character for digit.
    if (x == 0) { break; } // Stop if all digits written.
}
fwrite(&s[i], sizeof(char), 10-i, f); // Write ASCII bytes.
fclose(f);               // Close output file.
```

## Reading bytes

```
size_t fread(void *ptr ,  
            size_t size ,  
            size_t nmemb,  
            FILE *stream );
```

ptr: where to put the data we read.

size: the size of each data element in bytes.

nmemb: the number of data elements.

stream: the target file (opened with fopen()).

Very similar to fwrite()!

## Reading all the bytes in a file

```
#include <stdio.h>
#include <assert.h>

int main(int argc, char* argv[]) {
    FILE *f = fopen(argv[1], "r");
    unsigned char c;
    while (fread(&c, sizeof(char), 1, f) == 1) {
        printf("%3d_", (int)c);
        if (c > 31 && c < 127) {
            fwrite(&c, sizeof(char), 1, stdout);
        }
        printf("\n");
    }
}
```

## Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

# Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

```
$ ./fread-bytes fread-bytes.c
```

```
35 #
```

```
105 i
```

```
110 n
```

```
99 c
```

```
108 l
```

```
117 u
```

```
100 d
```

```
101 e
```

```
32
```

```
60 <
```

```
...
```

# Running fread-bytes

```
$ gcc -o fread-bytes -Wall -Wextra -pedantic fread-bytes.c
```

```
$ ./fread-bytes fread-bytes.c $ ./fread-bytes fread-bytes
```

35	#	127
105	i	69 E
110	n	76 L
99	c	70 F
108	l	2
117	u	1
100	d	1
101	e	0
32		0
60	<	0
...		...

## Text files versus binary files

- **To the system there is no difference between “text files” and “binary files”!**
- All files are just byte sequences.
- *Colloquially*: a text file is a file that is understandable when the bytes are interpreted as characters (in ASCII or some other character set).



# Text files versus binary files

- **To the system there is no difference between “text files” and “binary files”!**
- All files are just byte sequences.
- *Colloquially*: a text file is a file that is understandable when the bytes are interpreted as characters (in ASCII or some other character set).

## Compactness of storage

- A 32-bit integer takes up to 12 bytes to store as base-10 ASCII digits
- 4 bytes as raw data
- **Raw data takes up less space and is much faster to read.**
- But we need special programs to decode the data to human-readable form.

# IO summary

- Use `printf()` for text output.
- (And `scanf()` for text *input*.)
- Use `fwrite()` to write raw data.
- Use `fread()` to read raw data.
- Raw data files are more compact and faster to read/write.